

Per-request latency distribution, baseline vs. salted

